

Lecture 11. Neural Network Fundamentals

COMP90051 Statistical Machine Learning

Lecturer: Zahra Dasht Bozorgi







THE UNIVERSITY OF
MELBOURNE

Limitations of linear models

Some problems are linearly separable, but many are not

- In/out value 1
- In/out value 0



Possible solution: **composition**

$$x_1 \text{ XOR } x_2 = (x_1 \text{ OR } x_2) \text{ AND not}(x_1 \text{ AND } x_2)$$

We are going to compose perceptrons ...

Perceptron is *sort of* a building block for ANN

- ANNs are not restricted to binary classification
- Nodes in ANN can have various **activation functions**

Step function

$$f(s) = \begin{cases} 1, & \text{if } s \geq 0 \\ 0, & \text{if } s < 0 \end{cases}$$

Sign function

$$f(s) = \begin{cases} 1, & \text{if } s \geq 0 \\ -1, & \text{if } s < 0 \end{cases}$$

Logistic function


$$f(s) = \frac{1}{1 + e^{-s}}$$

tanh function

$$f(s) = \frac{e^s - e^{-s}}{e^s + e^{-s}}$$

Rectified linear unit

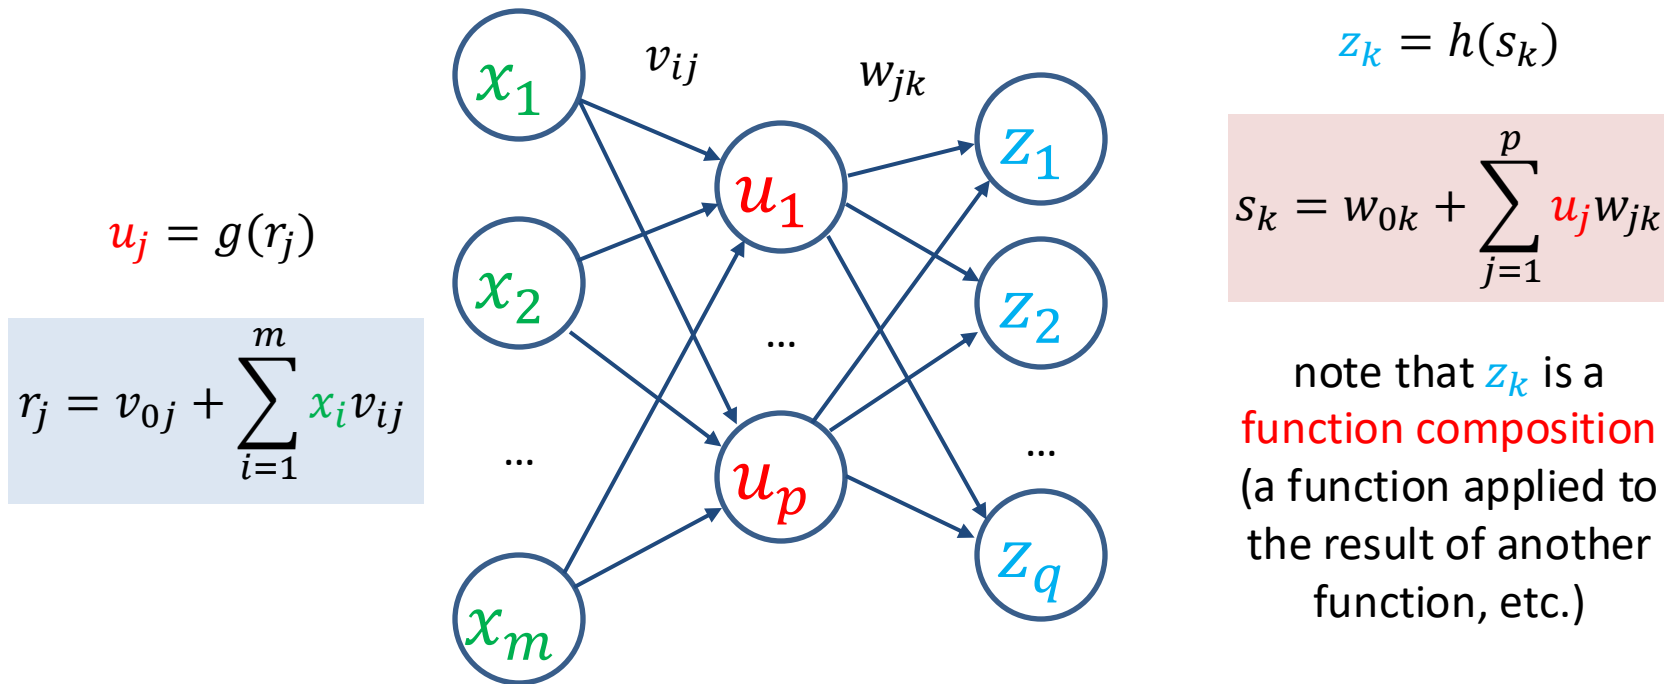
$$f(s) = \max\{0, s\}$$



[This Photo](#) by Unknown Author is licensed under [CC BY-SA](#)

...many others, many variations...

ANN as function composition



note that z_k is a **function composition** (a function applied to the result of another function, etc.)



here g, h are activation functions. These can be either same (e.g., both sigmoid) or different

you can add **bias node** $x_0 = 1$ to simplify equations: $r_j = \sum_{i=0}^m x_i v_{ij}$

similarly you can add bias node $u_0 = 1$ to simplify equations: $s_k = \sum_{j=0}^p u_j w_{jk}$


The power of ANN as a non-linear model

- ANNs are capable of approximating plethora non-linear functions, e.g., $z(x) = x^2$ and $z(x) = \sin x$
- For example, consider the following network. In this example, hidden unit activation functions are tanh



The power of ANN as a non-linear model

- ANNs are capable of approximating various non-linear functions, e.g., $z(x) = x^2$ and $z(x) = \sin x$



Blue points are the function values evaluated at different x . Red lines are the predictions from the ANN. Dashed lines are outputs of the hidden units

Adapted from Bishop 2007

- Universal approximation theorem** (Cybenko 1989): An ANN with a hidden layer with a finite number of units, and mild assumptions on the activation function, can approximate continuous functions on \mathbb{R}^n arbitrarily well

Representational capacity

- ANNs with a single hidden layer are **universal approximators**
- For example, such ANNs can represent any Boolean function

$$OR(x_1, x_2) \quad u = g(x_1 + x_2 - 0.5)$$

$$AND(x_1, x_2) \quad u = g(x_1 + x_2 - 1.5)$$


$$NOT(x_1) \quad u = g(-x_1)$$

$$g(r) = 1 \text{ if } r \geq 0 \text{ and } g(r) = 0 \text{ otherwise}$$

- Any Boolean function over m variables can be implemented using a hidden layer with up to 2^m elements
- More **efficient to stack** several hidden layers

“Depth” refers
to number of
hidden layers

Deep networks



$$\mathbf{s} = \tanh(\mathbf{A}'\mathbf{x}) \quad \mathbf{t} = \tanh(\mathbf{B}'\mathbf{s}) \quad \mathbf{u} = \tanh(\mathbf{C}'\mathbf{t}) \quad \mathbf{z} = \tanh(\mathbf{D}'\mathbf{u})$$

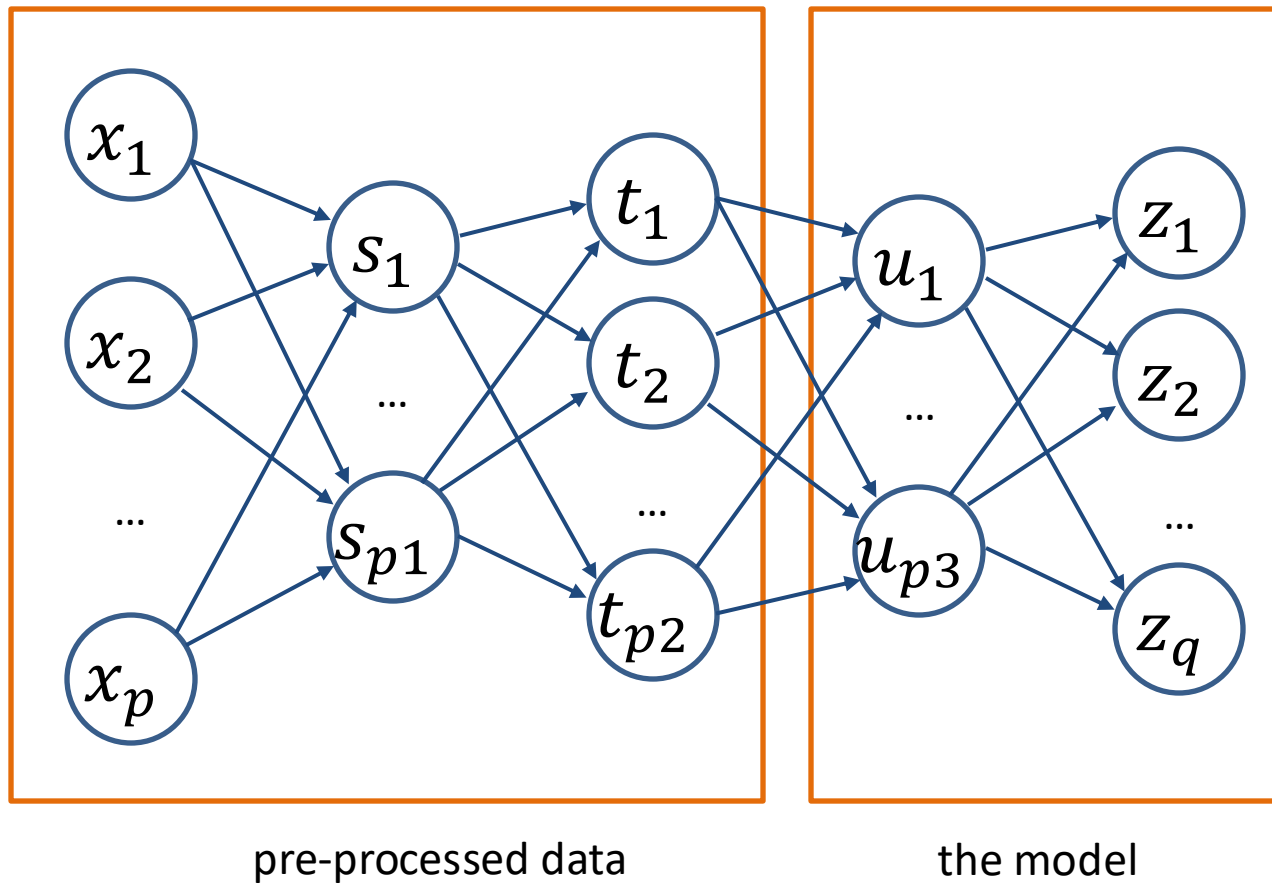
Deep ANNs as representation learning

- Consecutive layers form **representations** of the input of increasing complexity
- An ANN can have a simple *linear* output layer, but using complex *non-linear* representation

$$\mathbf{z} = \tanh \left(\mathbf{D}' \left(\tanh \left(\mathbf{C}' \left(\tanh \left(\mathbf{B}' \left(\tanh \left(\mathbf{A}' \mathbf{x} \right) \right) \right) \right) \right) \right) \right)$$

- Equivalently, a hidden layer can be thought of as the transformed feature space, e.g., $\mathbf{u} = \varphi(\mathbf{x})$
compare to **basis expansion** / **kernel learning**
- Parameters of such a transformation are learned from data

ANN layers as data transformation



Depth vs width

- A single arbitrarily wide layer in theory gives a universal approximator
- However (empirically) depth yields more accurate models
Biological inspiration from the eye:
 - * first detect small edges and color patches;
 - * compose these into smaller shapes;
 - * building to more complex detectors, of e.g. textures, faces, etc.
- Seek to mimic layered complexity in a network
- However **vanishing gradient problem** affects learning with very deep models
- Training error decreases but then increases and converges, vanishing gradient and large learning rate

Recap: Stochastic gradient descent training

Choose initial guess $\theta^{(0)}$, $k = 0$

Here θ is a set of all weights form all layers

For i from 1 to T (epochs)

For j from 1 to N (training examples)

Consider example $\{x_j, y_j\}$

Update: $\theta^{(k+1)} = \theta^{(k)} - \eta \nabla L(\theta^{(k)}); k \leftarrow k+1$

For regression we might use


$$L = \frac{1}{2} (z_j - y_j)^2$$

Need to compute partial derivatives $\frac{\partial L}{\partial v_{ij}}$ and $\frac{\partial L}{\partial w_j}$

Backpropagation: start with the chain rule

- Recall that the output z of an ANN is a function composition, and hence $L(z)$ is also a composition
 - $L = 0.5(z - y)^2 = 0.5(h(s) - y)^2 = 0.5(s - y)^2$
 - $= 0.5\left(\sum_{j=0}^p u_j w_j - y\right)^2 = 0.5\left(\sum_{j=0}^p g(r_j) w_j - y\right)^2 = \dots$
- Backpropagation makes use of this fact by applying the **chain rule** for derivatives

- $$\frac{\partial L}{\partial w_j} = \frac{\partial L}{\partial z} \frac{\partial z}{\partial s} \frac{\partial s}{\partial w_j}$$
- $$\frac{\partial L}{\partial v_{ij}} = \frac{\partial L}{\partial z} \frac{\partial z}{\partial s} \frac{\partial s}{\partial u_j} \frac{\partial u_j}{\partial r_j} \frac{\partial r_j}{\partial v_{ij}}$$



- Recall that


- $s = \sum_{j=0}^p u_j w_j$
- $r_j = \sum_{i=0}^m x_i v_{ij}$

Backpropagation: intermediate step

- Apply the chain rule

- $$\frac{\partial L}{\partial w_j} = \frac{\partial L}{\partial z} \frac{\partial z}{\partial s} \frac{\partial s}{\partial w_j}$$

- $$\frac{\partial L}{\partial v_{ij}} = \frac{\partial L}{\partial z} \frac{\partial z}{\partial s} \frac{\partial s}{\partial u_j} \frac{\partial u_j}{\partial r_j} \frac{\partial r_j}{\partial v_{ij}}$$



- Now define

$$\delta \equiv \frac{\partial L}{\partial s} = \frac{\partial L}{\partial z} \frac{\partial z}{\partial s}$$

$$\varepsilon_j \equiv \frac{\partial L}{\partial r_j} = \frac{\partial L}{\partial z} \frac{\partial z}{\partial s} \frac{\partial s}{\partial u_j} \frac{\partial u_j}{\partial r_j}$$

- Here $L = 0.5(z - y)^2$ and $z = s$

Thus $\delta = (z - y)$

- Here $s = \sum_{j=0}^p u_j w_j$ and $u_j = g(r_j)$

Thus $\varepsilon_j = \delta w_j g'(r_j)$

Backpropagation equations

- We have

$$* \frac{\partial L}{\partial w_j} = \delta \frac{\partial s}{\partial w_j}$$

$$* \frac{\partial L}{\partial v_{ij}} = \varepsilon_j \frac{\partial r_j}{\partial v_{ij}}$$

- ... where

$$* \delta = \frac{\partial L}{\partial s} = (z - y)$$

$$* \varepsilon_j = \frac{\partial L}{\partial r_j} = \delta w_j g'(r_j)$$

- Recall that

$$* s = \sum_{j=0}^p u_j w_j$$


$$* r_j = \sum_{i=0}^m x_i v_{ij}$$

- So $\frac{\partial s}{\partial w_j} = u_j$ and $\frac{\partial r_j}{\partial v_{ij}} = x_i$

- We have

$$* \frac{\partial L}{\partial w_j} = \delta u_j = (z - y) u_j$$

$$* \frac{\partial L}{\partial v_{ij}} = \varepsilon_j x_i = \delta w_j g'(r_j) x_i$$




Modern DNN libraries have active derivatives pre-computed; autodiff computes derivatives efficiently

Forward propagation

- Use current estimates of v_{ij} and w_j



- Calculate r_j , u_j , s and z



- Backpropagation equations


$$* \frac{\partial L}{\partial w_j} = \delta u_j = (z - y) u_j$$

$$* \frac{\partial L}{\partial v_{ij}} = \varepsilon_j x_i = \delta w_j g'(r_j) x_i$$

Backward propagation of errors

$$\frac{\partial L}{\partial v_{ij}} = \varepsilon_j x_i \leftarrow \boxed{\varepsilon_j = \delta w_j g'(r_j)} \leftarrow \frac{\partial L}{\partial w_j} = \delta u_j \leftarrow \boxed{\delta = (z - y)}$$

Notice how intermediate values get **reused**.



- Backpropagation equations

- * $\frac{\partial L}{\partial w_j} = \delta u_j = \boxed{(z - y)} u_j$

- * $\frac{\partial L}{\partial v_{ij}} = \varepsilon_j x_i = \boxed{\delta w_j g'(r_j)} x_i$