

Lecture 14. Recurrent Neural Networks, Attention, and the Transformer

COMP90051 Statistical Machine Learning

Zahra Dasht Bozorgi



THE UNIVERSITY OF
MELBOURNE

This lecture

- Recurrent networks for modelling sequences
 - * recurrent units
 - * back-propagation through time
 - * long-short term memory
- Transformers and attention

Recurrent Networks

*A DNN tailored to variable length
sequential inputs*

Sequential input

- Until now, we have assumed fixed-sized input
 - * Vectors of features \mathbf{x} in d dimensions
 - * Matrices of pixels in an image
- What if our input is a **sequence**?
 - * Frames in a video clip
 - * Time steps in an audio clip
 - * Words in a sentence
 - * A protein sequence
 - * Stock prices over time ...
- How can we model this in a DNN?

FCNNs poor for sequences


- Consider classifying sentences
 - * “This is the worst movie of all time, a real stinker” → 😞
 - * “The movie is a real stinker” → 😞
- Issue: inputs are *different lengths*
 - * **pad** them with empty “words” to be a fixed size
- Issue: how do we *represent words* as vectors?
 - * learn an “**embedding**” vector for each word
- Issue: phrases have *similar meaning even when at different locations*
 - * “a real stinker” is a key predictive feature
 - * if we naively apply FCNN needs to learn this concept repeatedly

ConvNets for Sequences?

- Sequences are just rectangular shaped images (e.g., embedding dim. times length): apply CNNs
 - * With **1D filters**
 - * The filter parameters are shared across time, and can find patterns in the input
- This is called the *time delay neural network*
- Downside:
 - * receptive field of filters are limited to finite size, i.e., the width of the convolutional filters, which can be expanded with deeper networks


Recurrent Neural Nets (RNNs)

- RNNs create networks dynamically, based on input sequence
 - * given sequence of inputs $\mathbf{x}^{(1)}, \mathbf{x}^{(2)}, \dots, \mathbf{x}^{(t)}$
 - * process each symbol from left to right, to form a sequence of hidden states $\mathbf{h}^{(t)}$
 - * each $\mathbf{h}^{(t)}$ encodes all inputs up to t



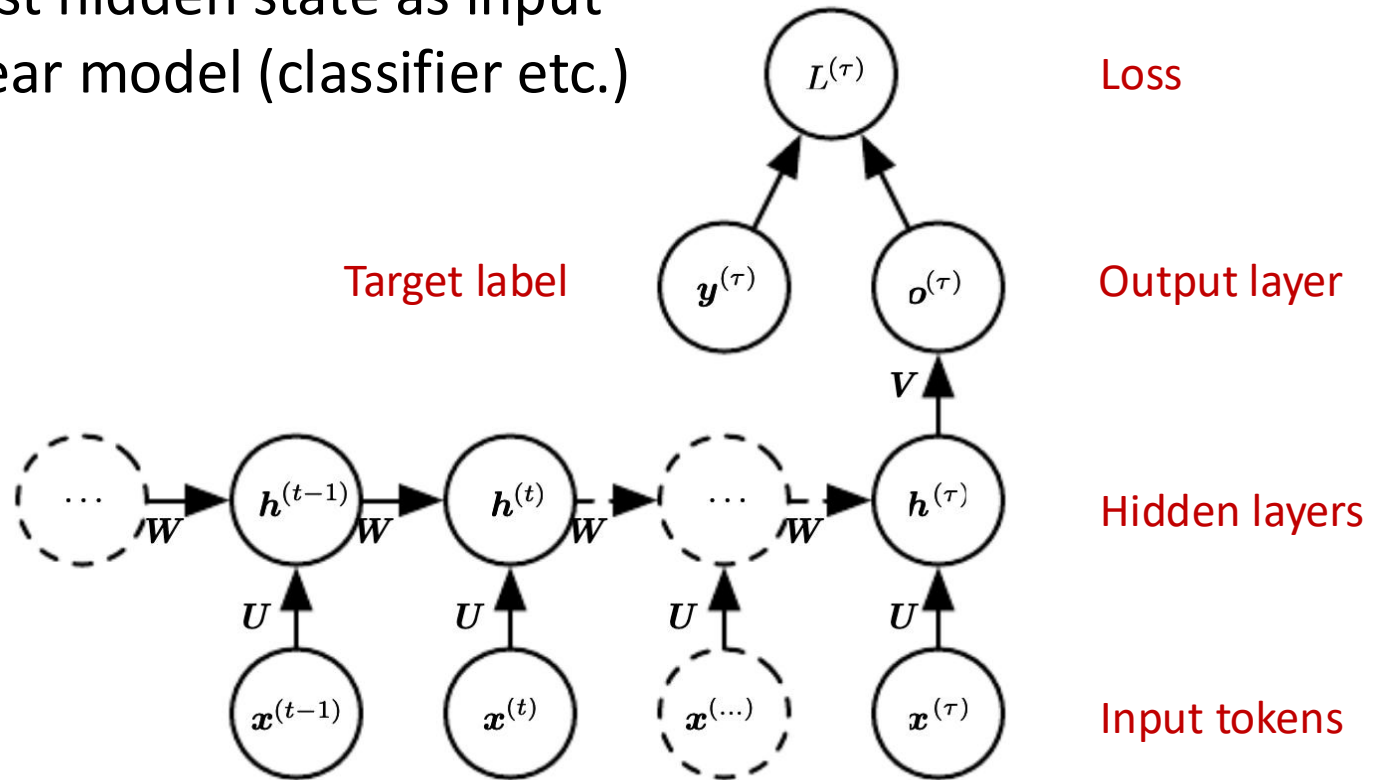
RNNs as Very Deep Networks

- Compared to NNets we've seen before:
 - * unfolded RNN has **depth equal to input sequence length**
 - * **parameters shared** between layers
- Can easily be 'unrolled' to cater to any input length



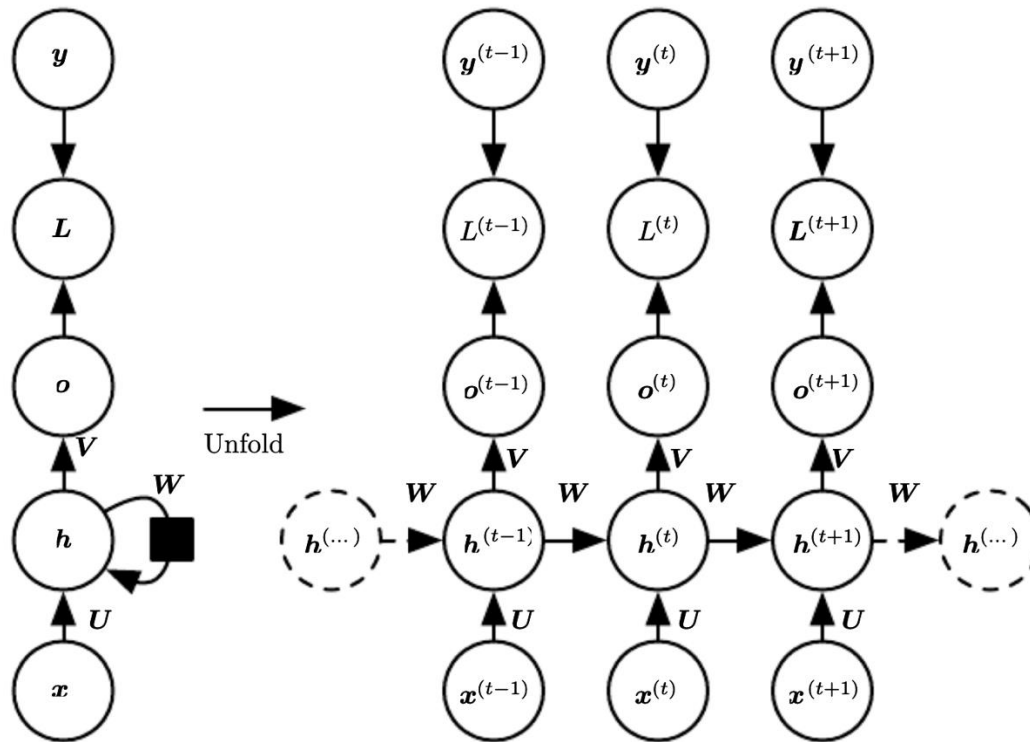
RNN Applications: Seq. Classification

- Sequence classification: labelling sequence
 - * use last hidden state as input to linear model (classifier etc.)




Sequence Tagging RNN

- Assign each item/token a label in sequence
 - * Given targets per item, can measure loss per item



Encoder-Decoder for Sequence Translation



E.g., English to French

Encoder RNN encodes input sequence into a context

Decoder RNN acts like a tagger, where we're trying to (re)generate next inputs

RNN Parameterisation

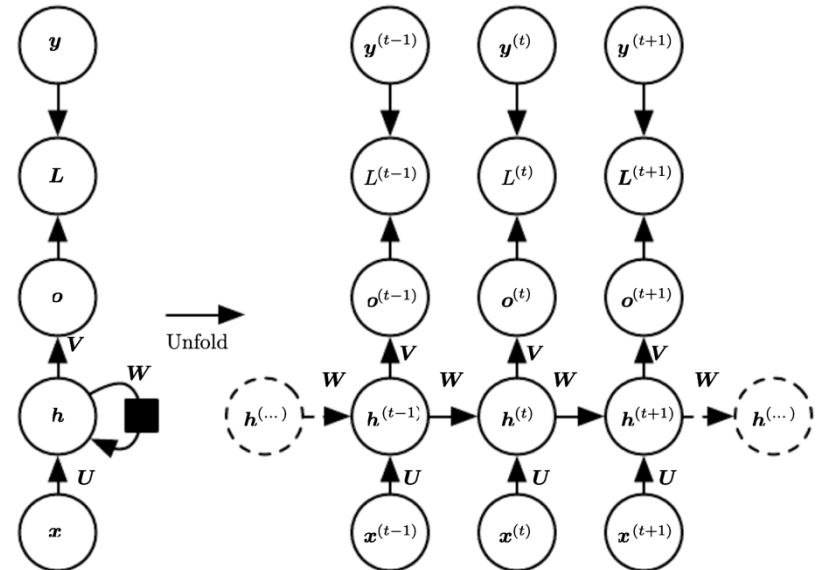
- Consider tagging RNN:
define f as follows

$$\mathbf{a}^{(t)} = \mathbf{b} + \mathbf{W}\mathbf{h}^{(t-1)} + \mathbf{U}\mathbf{x}^{(t)},$$

$$\mathbf{h}^{(t)} = \tanh(\mathbf{a}^{(t)}),$$

$$\mathbf{o}^{(t)} = \mathbf{c} + \mathbf{V}\mathbf{h}^{(t)},$$

$$\hat{\mathbf{y}}^{(t)} = \text{softmax}(\mathbf{o}^{(t)}),$$




- Parameters are $\mathbf{b}, \mathbf{W}, \mathbf{U}, \mathbf{c}, \mathbf{V}$
 - * not specific to timestep t , but shared across all positions
 - * this “template” can get unrolled arbitrarily

Training RNNs: Backprop. Thru. Time

- Backpropagation algorithms can be applied to network
 - * Called **backpropagation through time (BPTT)**
 - * Gradients from the loss at every position must be propagated back to the very start of the network
- Suffers from **gradient vanishing** problem
 - * Consider linear RNN, gradients of $\frac{\partial \mathbf{h}^{(T)}}{\partial \mathbf{h}^{(1)}} = \mathbf{W}^{T-1}$, thus can explode or vanish with large T , depending on largest eigenvalue of \mathbf{W} (i.e., greater than / less than one).
 - * Can't *learn* long distance phenomena (over 10+ steps)

Long Short-Term Memory (LSTM)

- In RNN, previous state is provided as an input
 - * Multiplied by weight matrix, and non-linearity applied
- LSTM introduces state self-loop, based on copying
 - * Takes **copy** of previous state, scaled by sigmoid **forget gate**
- Gradient magnitude now maintained
 - * Can handle 100+ distance phenomena (vs 5-10 for RNN)



Mini-summary

- Recurrent networks for modelling sequences
 - * recurrent units
 - * back-propagation through time
 - * long-short term memory
 - * applications

next: Transformers


Transformers

~~*More than meets the eye.*~~

*A method for processing sequence inputs in highly parallelizable manner, using **attention**.*



Attention

- **RNNs** over long sequences not too good at representing properties of the full sequence
 - * **Biased** towards the end (or ends) of the sequence
 - * Last hidden layer / context: A **bottleneck**!
- **Attention** averages over hidden sequence
 - * $\mathbf{c} = \sum_j \alpha_j \mathbf{h}^{(j)}$ summary weighted average
 - * $\alpha_j = \exp(e_j) / (\sum_{j'} \exp(e_{j'}))$ softmax
 - * $e_j = f(\mathbf{h}^{(j)})$ focuses at
- E.g., key phrase in review



Repeated attention in Seq2seq models

- Consider multiple sequential outputs, $\mathbf{s}^{(i)}$
 - * $\mathbf{c}_i = \sum_j \alpha_{ij} \mathbf{h}^{(j)}$
 - * $\mathbf{s}^{(i)} = f(\mathbf{s}^{(i-1)}, y^{(i-1)}, \mathbf{c}_i)$
 - * $\alpha_{ij} = \exp(e_{ij}) / (\sum_{j'} \exp(e_{ij'}))$
 - * $e_{ij} = f(\mathbf{s}^{(i-1)}, \mathbf{h}^{(j)})$
- Avoids bottleneck, and uncovers meaningful structure


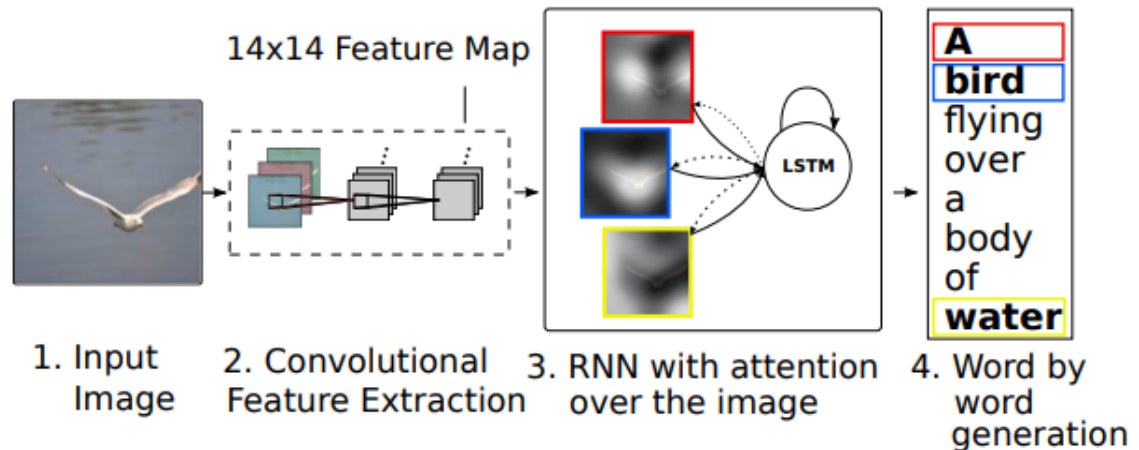


Neural Machine Translation by Jointly Learning to Align and Translate.

Bahdanau, Cho, Bengio, ICLR 2015

Attention in Vision

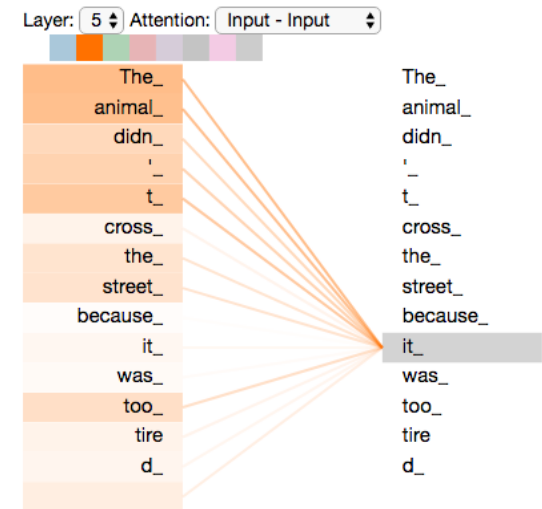
- Can attend to other representations, e.g., images
 - * Attention over matrix input
 - * RNNs during generation of caption



Show, Attend and Tell: Neural Image Caption Generation with Visual Attention,
Xu et al, ICML 2015

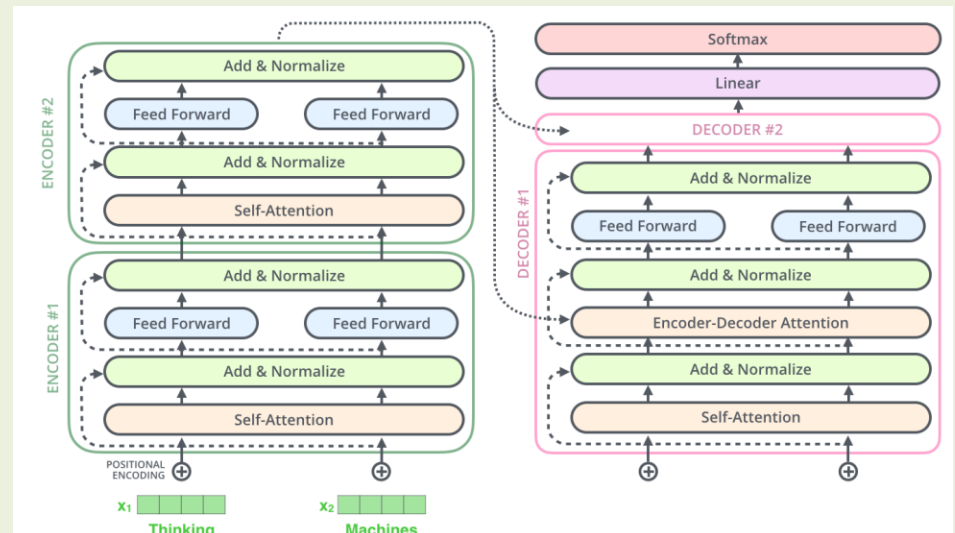
Self-attention

- **Transformers** use attention as means of representing sequences directly, instead of RNN
 - * Representation of item i is based on attention to the rest of the sequence
 - * Use item i as the query in attention against all items $j \neq i$
- Compared to RNNs
 - * No explicit position information (add to each symbol position index)
 - * Cheap: easily done in parallel



Transformer

- The Transformer uses self-attention as its main step
 - * Alongside residual, and normalization layers
 - * Using multiple “attention heads”, and deep stacking
- Applied first to translation
 - * Then raw text, e.g., BERT, RoBERTa, GPT
 - * Highly scalable
 - * Large performance gains over RNN models



This lecture

- Recurrent networks for modelling sequences
 - * recurrent units
 - * back-propagation through time
 - * long-short term memory
- Transformers and attention

next: Cross-validation