

# Lecture 16. Learning with expert advice

COMP90051 Statistical Machine Learning

Lecturer: Zahra Dasht Bozorgi



THE UNIVERSITY OF  
MELBOURNE

# This lecture


- Learning from expert advice / multiplicative weights
  - \* Learner listens to some/all experts making predictions
  - \* True outcomes are ADVERSARIAL!
  - \* Learner updates weights over experts based on their losses
  - \* Algorithms all forms of “multiplicative weights”
  - \* Nice clean bounds on total mistakes/loss: by “potential function” technique
- Infallible expert (one always perfect)
  - \* Majority Vote Algorithm
- Imperfect experts (none guaranteed perfect) – increasingly better
  - \* Weighted Majority Vote Algorithm by Halving
  - \* Weighted Majority Voting by General Multiplicative Weights
  - \* Probabilistic Experts Algorithm

# An infallible expert and the Majority Algorithm

*Warming up example*

# Warm-up: Case of the infallible expert

- **Experts**  $E_1, \dots, E_n$  predict the stock market daily
  - \* Each expert prediction is binary: stocks will go up/down
- Learner's game, daily:
  - \* Observe predictions of all experts
  - \* Make a prediction of its own
  - \* Observe outcome (could be anything!)
  - \* Goal: minimise number total mistakes
- **Infallible expert** assumption:
  - \* 1 or more experts makes no mistakes



# Infalible Expert Algorithm: Majority Vote

1. Initialise set of experts who haven't made mistakes  $E = \{1, \dots, n\}$
2. Repeat per round
  - a) Observe predictions  $E_i$  for all  $i \in \{1, \dots, n\}$
  - b) Make **majority prediction**  
 $\arg \max_{y \in \{-1, 1\}} \sum_{i \in E} 1[E_i = y]$
  - c) Observe correct outcome
  - d) **Remove mistaken experts** from  $E$



CCA3.0: Krisada, Noun Project

# Mistake Bound for Majority Vote

Proposition: Under infallible expert assumption, majority vote makes total mistakes  $M \leq \log_2 n$

Intuition: Halving  
(e.g. tree data  
structures!)? Expect  
to see log

## Proof

- Loop invariant: If algorithm makes a mistake, then at least  $|E|/2$  experts must have been wrong
- I.e. for every incorrect prediction,  $E$  reduced by at least half. I.e. after  $M$  mistakes,  $|E| \leq n/2^M$
- By infallibility, at all times  $1 \leq |E|$
- Combine to  $1 \leq |E| \leq n/2^M$ , then solve for  $M$ .

**$|E|$  is the  
potential  
function**


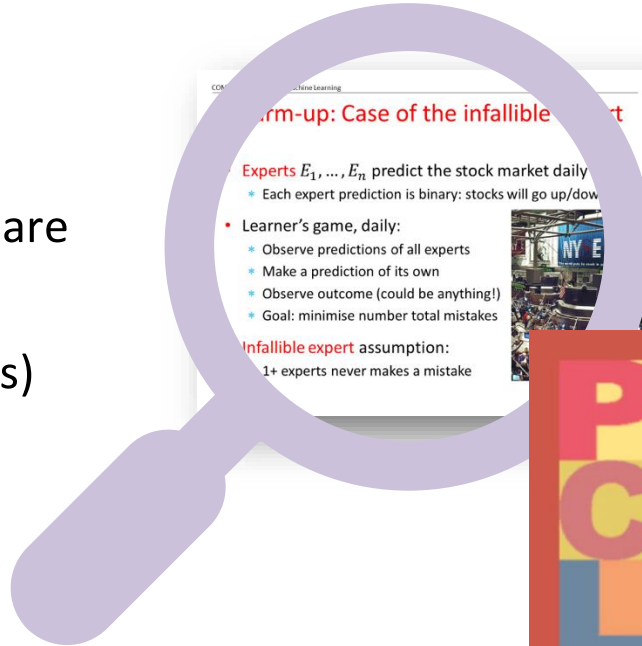
# How is this “online learning”?

## Learning

- **Weights on which experts** are worth listening to
- (Infallible case: 0/1 weights)
- Making predictions/  
taking actions
- Incurring loss (so far 0/1)
- IID “Distribution” replaced by **adversarial outcomes**

## Online

- A **repeated game**



# Mini Summary

- Learning with expert advice paradigm
  - \* Abstraction of online learning problem
  - \* Adversarial feedback
  - \* Later: Applications abound
- Bounds on mistakes (later losses) “easy”
  - \* Involve “potential function” technique
  - \* Later: interested in scaling with best expert performance

Next: Imperfect experts. Don't drop bad experts



# Imperfect experts and the Halving Algorithm

*Similar proof technique; similar algorithm;  
much more interesting setting*

# No one's perfect

- No more guarantee of an infallible expert
- What breaks?
  - \* We could end up with  $E = \emptyset$ , how to predict then?
  - \* No sense: “Zero tolerance” dropping experts on a mistake
- Very general setting / very few assumptions
  - \* Not assuming anything about expert error rates
  - \* Not assuming anything about correlation of expert errors
  - \* Not assuming anything about outcome observations. Not even stochastic (could be adversarial!)

# Imperfect experts: Halving Algorithm

1. Initialise  $w_i = 1$  weight of expert  $E_i$
2. Repeat per round
  - a) Observe predictions  $E_i$   
for all  $i \in \{1, \dots, n\}$
  - b) Make **weighted majority prediction**  
 $\arg \max_{y \in \{-1, 1\}} \sum_{i \in E} w_i 1[E_i = y]$
  - c) Observe correct outcome
  - d) **Downweigh each mistaken expert  $E_i$**   
 $w_i \leftarrow w_i / 2$



CCA3.0: Krisada, Noun Project

# Mistake Bound for Halving

Proposition: If the best expert makes  $m$  mistakes, then weighted majority vote makes  $M \leq 2.4(m + \log_2 n)$  mistakes.

## Proof

- Invariant: If algorithm makes a mistake, then weight of wrong experts is at least half the total weight  $W = \sum_{i=1}^n w_i$
- Weight of wrong experts reduced by  $1/2$ , therefore total weight reduced by at least  $3/4$ . I.e. after  $M$  mistakes,  $W \leq n(3/4)^M$
- Best expert  $E_i$  has  $w_i = (1/2)^m$
- Combine to  $(1/2)^m = w_i \leq W \leq n(3/4)^M$
- Taking logs  $-m \leq \log_2 n + M \log_2(3/4)$ , solving  $M \leq \frac{m + \log_2 n}{\log_2(4/3)}$

# Compare, compare: What's going on?

- Price of imperfection (vs. infallibility) is  $\mathcal{O}(m)$ 
  - \* Infallible case:  $M \in \mathcal{O}(\log n)$
  - \* Imperfect case:  $M \in \mathcal{O}(m + \log n)$
- Scaling to many experts is no problem
- Online learning vs. PAC frameworks

	Modelling of losses	Ultimate goal
<b>PAC</b>	i.i.d. losses (due to e.g. Hoeffding)	(For ERM; L6c) Small estimation error $R[f_m] - R[f^*]$ . Bounded in terms of family's VC dimension
<b>Online learning</b>	Adversarial/arbitrary losses	Small $M - m$ . Bounded in terms of number of experts.

# Mini Summary

- Imperfect expert setting
  - \* Don't drop bad experts, just halve their weight
  - \* Predict by weighted majority, not simply majority
  - \* Mistake bound follows similar “potential function” pattern!
- Learning with expert advice paradigm
  - \* Key difference to PAC is adversarial feedback
  - \* Similarity: Also concerned with performance relative to “best in class”

Next: Imperfect experts continued. Generalising halving.

# From Halving to Multiplying weights by $1 - \varepsilon$

*Generalising weighted majority.*

# Useful (but otherwise boring) inequalities

- Lemma 1: For any  $\varepsilon \in [0, 0.5]$ , we have
$$-\varepsilon - \varepsilon^2 \leq \log_e(1 - \varepsilon) < -\varepsilon$$

- Proof:
  - \* Upper bound by Taylor expansion, dropping all but first term (as they're negative)
  - \* Lower bound by convexity of  $\exp(-\varepsilon - \varepsilon^2)$


- Lemma 2: For all  $\varepsilon \in [0, 1]$  we have,
$$1 - \varepsilon x > \begin{cases} (1 - \varepsilon)^x, & \text{if } x \in [0, 1] \\ (1 + \varepsilon)^{-x}, & \text{if } x \in [-1, 0] \end{cases}$$

- Proof: by convexity of the RHS functions



# Weighted Majority Vote Algorithm

1. Initialise  $w_i = 1$  weight of expert  $E_i$
2. Repeat per round
  - a) Observe predictions  $E_i$  for all  $i \in \{1, \dots, n\}$
  - b) Make weighted majority prediction  $\arg \max_{y \in \{-1, 1\}} \sum_{i \in E} w_i 1[E_i = y]$
  - c) Observe correct outcome
  - d) Downweigh each mistaken expert  $E_i$   
 $w_i \leftarrow (1 - \varepsilon)w_i$



CCA3.0: Krisada, Noun Project

# Mistake Bound

Proposition: If the best expert makes  $m$  mistakes, then  $(1 - \varepsilon)$ -weighted majority makes  $M \leq 2(1 + \varepsilon)m + (2\log_e n)/\varepsilon$  mistakes.

Bound improves dependence on  $m$  compared to halving. **Why?**

## Proof

- Whenever learner mistakes, at least half of total weight reduced by factor of  $1 - \varepsilon$ . So after  $M$  mistakes,  $W \leq n(1 - \varepsilon/2)^M$
- Best expert  $E_i$  has  $w_i = (1 - \varepsilon)^m$
- Combine to  $(1 - \varepsilon)^m = w_i \leq W \leq n(1 - \varepsilon/2)^M$
- Taking logs:  $m\log_e(1 - \varepsilon) \leq \log_e n + M\log_e(1 - \varepsilon/2)$
- Lemma 1 replaces both  $\log_e(1 - \varepsilon)$ :  $-m(\varepsilon + \varepsilon^2) \leq \log_e n - M\varepsilon/2$
- Solving for  $M$  proves the bound.

# Dependence in $m$ provably near optimal!

- New to lower bounds? example shows an analysis or even an algorithm can't do better than some limit
- Weighted majority almost achieves  $2m$  dependence, with  $2(1 + \varepsilon)m$  (considering no. experts fixed)
- Example with  $M = 2m$ 
  - \* Consider  $n = 2$  with  $E_1$  ( $E_2$ ) correct on odd (even) days
  - \* Then best expert makes mistakes half the time
  - \* But after 1<sup>st</sup> round, for any  $\varepsilon$ , majority vote is wrong all the time, as incorrect expert gets more than half weight
- Consequence? Can't improve the constant 2 factor in  $m$

# Mini Summary

- Imperfect expert setting continued...
- From halving to multiplicative weights!
  - \* Mistake bound proved as usual via “potential function” trick
  - \* Bound’s dependence on best expert improved to  $2 + \varepsilon$  factor
- Lower bound / impossibility result
  - \* Factor of 2 is optimal for (deterministic) multiplicative weights!

Next: Imperfect experts continued. Randomise!!

# The probabilistic experts algorithm


*wherein randomisation helps us do better!*

# Probabilistic experts algorithm

- Change 1 from mistakes: **Loss**  $\ell_i^{(t)} \in [0,1]$  of  $E_i$ , round  $t$
  - Change 2: Randomised algorithm means, bounding **expected losses** (sound familiar? It should.... a risk!)
1. Initialise  $w_i = 1$  weight of expert  $E_i$
  2. Repeat per round
    - a) Observe predictions  $E_i$  for all  $i \in \{1, \dots, n\}$
    - b) Predict  $E_i$  of expert  $i$  **with probability**  $\frac{w_i}{W}$  where  $W = \sum_{i=1}^n w_i$
    - c) Observe **losses**
    - d) Update each weight  $w_i \leftarrow (1 - \varepsilon)^{\ell_i^{(t)}} w_i$

# Probabilistic experts: Expected loss bound

- Proposition: Expected loss of the probabilistic experts algorithm is  $L \leq \frac{\log_e n}{\varepsilon} + (1 + \varepsilon)L^*$  where  $L^*$  is the minimum loss over experts.
- Proof: next, follows similar “potential” pattern
- Beats deterministic! Shaves off optimal constant 2
- Generalises in many directions. Active area of research in ML, control, economics, in top labs.



# Proof: Upper bounding potential function

- Learner's round  $t$  expected loss:  $L_t = \frac{\sum_{i=1}^n w_i^{(t)} \ell_i^{(t)}}{W(t)}$
- By Lemma 2, since losses are  $[0,1]$ :  
updated  $w_i^{(t+1)} \leftarrow (1 - \varepsilon)^{\ell_i^{(t)}} w_i^{(t)} \leq (1 - \varepsilon \ell_i^{(t)}) w_i^{(t)}$

- Rearrange to obtain recurrence relation:

$$\begin{aligned} W(t+1) &\leq \sum_{i=1}^n (1 - \varepsilon \ell_i^{(t)}) w_i^{(t)} = \sum_{i=1}^n w_i^{(t)} \left( 1 - \varepsilon \frac{\sum_{i=1}^n w_i^{(t)} \ell_i^{(t)}}{\sum_{i=1}^n w_i^{(t)}} \right) \\ &= W(t) (1 - \varepsilon L_t) \end{aligned}$$

- Initialisation gave  $W(0) = n$ , so telescoping we get:

$$W(T) \leq n \prod_{t=1}^T (1 - \varepsilon L_t)$$



# Proof: Lower bounding potential, Wrap up

- Proved upper bound:  $W(T) \leq n \prod_{t=1}^T (1 - \varepsilon L_t)$

- Lower bound from best expert total loss  $L^*$ :

$$W(T) \geq (1 - \varepsilon)^{L^*}$$

- Combining bounds and taking log's:

$$L^* \log_e(1 - \varepsilon) \leq \log_e n + \sum_{t=1}^T \log_e(1 - \varepsilon L_t)$$


- By Lemma 1:  $-L^*(\varepsilon + \varepsilon^2) \leq \log_e n - \varepsilon \sum_{t=1}^T L_t$

- Linearity of expectation  $L = \sum_{t=1}^T L_t$ , rearranging:

$$L \leq \frac{\log_e n}{\varepsilon} + (1 + \varepsilon)L^*$$

# Applications of multiplicative weights [Kale thesis 2007]

- Learning quantum states from noisy measurements
- Derandomising algorithms
- Solving certain zero-sum games
- Fast graph partitioning
- Fast solving of semidefinite programming problems
- Portfolio optimisation
- A basis for boosting
- Sparse vector technique in differential privacy



Public domain

# Mini Summary

- Introducing randomisation to learning with experts
  - \* Algorithm choosing a random expert to follow
  - \* Weights become probabilities
  - \* Mistakes generalise to losses
- Loss bound
  - \* Have to bound expected loss (hey, risk!!)
  - \* Shaves off that 2 factor. Proves that randomisation really helps!

Next: Only observe reward of chosen expert → bandits!